

Hidden Markov Models

- Probabilistic graphical model
- Generative Approach
- Fully connected: "ergodic" HMM ; "Left-Right" HMM:
no backlinks

Methods to work with HMMs

- 1.) Matching (FORWARD - / BACKWARD) Algorithm
 $\alpha_+(s_i)$ $\beta_+(s_i)$

$$\boxed{\text{find: } p(\langle o_1, \dots, o_n \mid \lambda \rangle)}$$

- 2.) Most likely sequence (of states) for given observation

$$\boxed{\operatorname{argmax}_{\langle s_1, \dots, s_n \rangle} p(\langle o_1, \dots, o_n \rangle, \langle s_1, \dots, s_n \rangle)}$$

Viterbi Algorithm

- 3.) Training : find $\lambda = (A, B, \pi)$ given training samples

Baum - Welch - Formulas

- FORWARD / BACKWARD
- EM - Algorithm

FORWARD Algorithm

Forward Variable:

$$\alpha_t(i) = P(\underbrace{o_1, \dots, o_t}_{\text{partial sequence observation}} \mid \underbrace{q_t = s_i}_{\text{state } s_i \text{ at time } t}, \lambda)$$

Initialization : $\alpha_1(i) = \pi_i b_i(o_1)$

Induction :
$$\alpha_{t+1}(i) = \left(\sum_{j=1}^N \alpha_t(j) \cdot a_{ji} \right) \cdot \underbrace{b_i(o_{t+1})}_{\text{observation probability of } o_{t+1} \text{ at state } i}$$

Termination :
$$P(o|\lambda) = \sum_{i=1}^N \alpha_T(i)$$

Computation requires N^2T calculations

BACKWARD Algorithm

Backward Variable:
$$\beta_t(i) = P(\underbrace{o_{t+1}, \dots, o_T}_{\text{partial sequence observation}} \mid \underbrace{q_t = s_i}_{\text{given state } i \text{ at time } t}, \lambda)$$

Initialization : $\beta_T(i) = 1$ (chosen arbitrarily)

Induction :
$$\beta_t(i) = \sum_{j=1}^N a_{ij} \cdot b_j(o_{t+1}) \cdot \beta_{t+1}(j)$$

Computation requires N^2T calculations

Solution for Task 2 (Most likely hidden state sequence)

Individually most likely states :

$$\gamma_+(i) = P(q_t = s_i | o, \lambda)$$

$$= \frac{\alpha_+(i) \cdot \beta_+(i)}{P(o | \lambda)} = \frac{\alpha_+(i) \beta_+(i)}{\sum_{i=1}^N \alpha_+(i) \beta_+(i)}$$

$$q_t = \underset{1 \leq i \leq N}{\operatorname{argmax}} \gamma_+(i)$$

(state with highest probability)

Problem: resulting state sequence might not even be possible ...

\Rightarrow Define best state sequence $P(Q | o, \lambda)$, which is equivalent to $P(Q, o | \lambda)$ maximization



Viterbi Algorithm

Same as FORWARD but:

- Replace every Σ (sum) by max
- Technical detail: store actual path in separate list ψ .

Solution for task 3: Training / Learning -

Determine $A, B, \vec{\pi}$: no analytical solution \Rightarrow solve iteratively

Baum-Welch-Formulas (FORWARD / BACKWARD + EM)

Expectation

$\xi_t(s_i, s_j)$: Probability for transition $s_i \rightarrow s_j$ at time t

$$\xi_t(s_i, s_j) = \frac{\alpha_t(s_i) \cdot a_{s_i s_j} \cdot b_{s_j}(o_{t+1}) \cdot \beta_{t+1}(s_j)}{\sum_{k=1}^M \sum_{l=1}^M \alpha_t(s_k) \cdot a_{s_k s_l} \cdot b_{s_l}(o_{t+1}) \cdot \beta_{t+1}(s_l)}$$

normalization for all transitions between time t and $t+1 \Rightarrow \xi$ is a PDF

$\gamma_t(s_i)$: Probability of being in state s_i at time t

$$\gamma_t(s_i) = \sum_{j=1}^M \xi_t(s_i, s_j)$$

$$\sum_{t=1}^T \xi_t(s_i, s_j) : E\#(s_i, s_j) \quad (s_i \rightarrow s_j)$$

$$\sum_{t=1}^T \gamma_t(s_i) : E\#(s_i) \quad (\text{from } s_i)$$

Maximization

$\bar{\pi}_i$: expected # of times in s_i at $t=1$

$$\bar{\pi}_i = \gamma_1(s_i)$$

$$a_{ij} = \frac{E\#(s_i, s_j)}{E\#(s_i)}$$

$$\bar{b}_{s_j}(k) = \frac{E\#(s_j) \text{ s.t. } o_t = k}{E\#(s_j)}$$

MARKOV RANDOM FIELD

$$[f_{ij}] = \underset{f_{ij}}{\operatorname{argmax}} P([g_{ij}] | [f_{ij}]) \cdot P([f_{ij}])$$

$$P(f_{ij} | f_{i1}, \dots, f_{iM}) = P(f_{ij} | N(f_{ij}))$$

Hammerley - Clifford Theorem : MRF \Leftrightarrow GRF

$$p(\vec{x}) = \frac{1}{Z} \cdot e^{-H(\vec{x})}$$

$$H(\vec{x}) = \sum_{m \in S} V_m(\vec{x})$$

with submodularity condition:

Define V_m s.t. similar vectors
yield low Energy values

$$E(0,0) + E(1,1) \leq E(0,1) + E(1,0)$$

For our image smoothing example:

Pairwise Potential ("regularizer")


$$E(f_{ij}, f' \in N(f_{ij})) = \|f_{ij} - f'\|_2^2$$

Unary Potential ("data term")


$$-\log(N(g_{ij}, f_{ij}, \sigma)) = \frac{(f_{ij} - g_{ij})^2}{\sigma}$$

Minimization Problem

$$\sum_{i=1}^N \left(E(x_i) + \sum_{j=1}^N E(x_i, x_j) \right)$$

Graph Cuts for Energy minimization

Edge

assignment 0 to x_{ij} assignment 1 to x_{ij}

[if included in cut]

Cut cost (for each edge cut):

Energy value unary / pairwise

\Rightarrow min cut finds x assignment for
unary and pairwise potentials

that leads to overall lowest Energy

 \Rightarrow highest $P(x_{ij})$ for the MRF \Rightarrow solution (for binary labels)L - Expansion Algorithm for L labels:

while (changes):

for each label l :binaryGraphCut(l) \Leftarrow 1 vs. all GraphCut's

"does the number of assignments
for label l grow?"

iterate until everything is
stable